

# Deep Learning Lab - Lecture 3

Last modified: September 30, 2024

---

Instructor: E. Vercesi

TAs: A. Dei Rossi, G. Dominici, S. Huber

vercee@usi.ch - {alvise.dei.rossi, gabriele.dominici,  
stefano.huber} @usi.ch

September 30, 2024

# Today plan/objective

- Linear regression with PyTorch
- Presentation of **Assignment 1**.
- You can start working on **Assignment 1** and / or ask questions on **Exercise 1** and / or ask questions on **Exercise 2**.

## Toy task - Linear regression i

- Let  $D, N \in \mathbb{N}_+$ .
  - $N$  data points
  - $D$  dimension of the input
- Hence, our dataset is


$$\mathcal{D} = \{(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^N, y^N) \mid \mathbf{x}^i \in \mathbb{R}^D, y^i \in \mathbb{R}\}$$

- Recall our goal: find  $f_{\theta}(\mathbf{x})$  that is a good approximation of the true function
- $\theta$  should be *learned* from data

## Toy task - Linear regression ii

- **Linear Regression:**  $f_{\theta}(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$ ,  $\mathbf{w} \in \mathbb{R}^D$ ,  $b \in \mathbb{R}$
- **Model parameters:**  $\theta = [\mathbf{w}, b]$
- Set  $D = 1$  in all that follows, namely we want to learn a “line” in the common sense

## Linear regression with $D = 1$



- Our linear model is  $f_{w,b}(x) = wx + b$  where  $w, b$  should be learnt from data

# Regression framework

This is a regression problem, we have to make some choices

- Loss function: Mean Squared Error

$$\text{MSE loss: } \mathcal{L}(w, b) := \frac{1}{N} \sum_{i=1}^N \|(wx^i + b) - y^i\|_2^2$$

- Optimizer: Stochastic Gradient Descent
- Hyperparameters:
  - (a) Learning rate  $\alpha$
  - (b) Epochs  $E$  number of times to see datasets
  - (c) Batch size  $B$  divide dataset into smaller

Note: you can compute the number of steps  $T$  as

$$T = \frac{EN}{B}$$

**Today:** no batches, no epochs, just steps!

# The algorithm

In this case, it is simple to design the algorithm step-by-step: after choosing the hyperparameters

- (1) Randomly initialize parameters  $w_0, b_0$
- (2) For each step  $t \leq T$ 
  - (i) Compute the loss  $\mathcal{L}(w_t, b_t)$
  - (ii) Compute the gradient

$$\left[ \frac{\partial \mathcal{L}}{\partial w}(w_t, b_t), \frac{\partial \mathcal{L}}{\partial b}(w_t, b_t) \right]$$

- (iii) Update parameters:

$$w_{t+1} = w_t - \alpha \frac{\partial \mathcal{L}}{\partial w}(w_t, b_t)$$

and

$$b_{t+1} = b_t - \alpha \frac{\partial \mathcal{L}}{\partial b}(w_t, b_t)$$

# Toy example

- Toy example means that:
  - We *pretend to know* the real function
  - We generate data from the true function *adding noise*
  - We end up with noisy data points (reality has noise)
- To generate data:
  - Start with the true  $w^*, b^*$  (we choose them)
  - Randomly sample  $x^i$
  - Add *Gaussian noise*, that is

$$y^i = w^* x^i + b^* + \varepsilon^i \quad \varepsilon^i \sim \mathcal{N}(0, \sigma)$$

- $(x^i, y^i)$  are our noisy data point



# A “create noisy data” function

```
1 def create_dataset(sample_size=10, sigma=0.1, w_star=1, b_star = 1,
2                     x_range=(-1, 1), seed=0):
3     # Set the random state in numpy
4     torch.manual_seed(seed)
5     # Unpack the values in x_range
6     x_min, x_max = x_range
7     # Sample sample_size points from a uniform distribution
8     X = torch.rand(sample_size)
9     # Get min and max
10    min_x_sampled = torch.min(X)
11    max_x_sampled = torch.max(X)
12    # Rescale between x_min and x_max
13    X = X * (x_max - x_min) + x_min
14    # Compute hat(y)
15    y_hat = X * w_star + b_star
16    # Compute y (Add Gaussian noise)
17    y = y_hat + torch.normal(torch.zeros(sample_size),
18                             sigma*torch.ones(sample_size))
19    return X, y
```

# Generate train & validation points


```
1 num_samples_train = 160
2 num_samples_validation = 20
3 num_samples_test = 20
4
5 # Set the seed
6 seed_train = 0
7 seed_validation = 1
8 seed_test = 2
9
10 # Set a value of noise (=sigma)
11 sigma = 1.3
12
13 # Define x_range
14 x = (-2, 2)
15
16 # Generate train data
17 X_train, y_train = create_dataset(
18     sample_size=num_samples_train, sigma=sigma, w_star=w_star,
19     b_star = b_star, x_range=x, seed=seed_train)
20
21 # Generate the validation data form the same distribution but with a different seed
22 X_val, y_val = create_dataset(
23     sample_size=num_samples_validation, sigma=sigma, w_star=w_star,
24     b_star = b_star, x_range=x, seed=seed_validation)
```

# Visualize the data <sub>i</sub>

Usually, plots are done with [Matplotlib](#) or [Seaborn](#)

```
1  fig, ax = plt.subplots(1, 3) # Create a subplot
2
3  # You can use torch directly with matplotlib
4  ax[0].plot(X_train, y_train, 'ko', alpha = 0.7)
5  ax[1].plot(X_val, y_val, 'ko', alpha = 0.7)
6  ax[2].plot(X_test, y_test, 'ko', alpha = 0.7)
7
8  # Just to have some values in the x-axis
9  x_range = torch.arange(start=min(x) - 0.1, end=max(x) + 0.1, step=0.01)
10 for i in range(3):
11     ax[i].set_ylim([-4, 9]) # You can see that these limits provide a better visualization
12     ax[i].plot(x_range, w_star * x_range + b_star)
```

## Visualize the data ii



# Linear regression

- We need a linear **model** → already implemented in PyTorch : `torch.nn.Linear`
- We need a Loss Function. Regression → Mean Square Error (MSE). `torch.nn.MSELoss`
- We need an optimizer, let's stick to stochastic gradient descent, already implemented in PyTorch: `torch.optim.SGD`
- Last but not least: we want to use GPUs!

```
1 DEVICE = torch.device('cuda' if torch.cuda.is_available() else 'mps')
2     if torch.backends.mps.is_available() else 'cpu')
```

# Linear regression

- We then create the model, defining the loss function and optimizer
- First hyperparameter to tune: learning rate

```
1 model = nn.Linear(1, 1) # Dimension of input: 1, dimension of output: 1  
2 loss_fn = nn.MSELoss()  
3 learning_rate = 0.5  
4 optimizer = optim.SGD(model.parameters(), lr=learning_rate)
```

# Linear regression

## What's happening at the beginning?

```
1 print("Initial w:", model.weight, "Initial b:\n", model.bias)
2 print("Value in x = 1:", model(torch.tensor([1.]))) # This computes  $w * 1 + b$ 
3 print("Actual value in x = 1:", w_star * 1 + b_star)
4 # Loss evaluation
5 initial_loss_function = loss_fn(X_train.reshape(-1, 1), y_train.reshape(-1, 1))
6 print("Initial loss function:", initial_loss_function)
```

```
1 Initial w: Parameter containing:
2 tensor([[ -0.2032]], requires_grad=True) Initial b:
3 Parameter containing:
4 tensor([ 0.5817]), requires_grad=True)
5 Value in x = 1: tensor([ 0.3785]), grad_fn=<ViewBackward0>)
6 Actual value in x = 1: 5
7 Initial loss function: tensor(10.7261)
```

We are almost ready for the training, but first of all, we have to put everything as needed for PyTorch

- Everything on GPU if we want to work on GPU
- Everything of the suitable type/shape

important put data in suitable dimensions

```
1 X_train = X_train.reshape(-1, 1).to(DEVICE)
2 y_train = y_train.reshape(-1, 1).to(DEVICE)
3 X_val = X_val.reshape(-1, 1).to(DEVICE)
4 y_val = y_val.reshape(-1, 1).to(DEVICE)
5 model = model.to(DEVICE)
```



# The training loop

```
1  n_steps = 10 # Number of updates of the gradient
2  for step in range(n_steps):
3      model.train() # Set the model in training mode
4      # Set the gradient to 0
5      optimizer.zero_grad() # Or model.zero_grad()
6      # Compute the output of the model
7      y_hat = model(X_train)
8      # Compute the loss
9      loss = loss_fn(y_hat, y_train)
10     # Compute the gradient
11     loss.backward()
12     # Update the parameters
13     optimizer.step()
```

gradients are accumulated, it's important to set `optimizer.zero_grad()` -> so important

# Evaluation process

Check what happens in the validation data

```
1  for step in range(n_steps):
2      # #####
3      # Code of previous slide here
4      # #####
5      with torch.no_grad(): #
6          # Compute the output of the model
7          y_hat_val = model(X_val)
8          # Compute the loss
9          loss_val = loss_fn(y_hat_val, y_val)
10         # At every step, print the losses
11         print("Step:", step, "- Loss eval:", loss_val.item())
```


loss function \infty, it's not a number.

# Training done!

Recall that we chose  $w^* = 3, b^* = 2$ .

```
1 print("Training done, with an evaluation loss of {}".format(loss_val.item()))
2 # Get the final value of the parameters
3 print("Final w:", model.weight, "Final b:\n", model.bias)
```

u should see  
something that decreases.



Training done, with an evaluation loss of 1.8968610763549805  
Final w: Parameter containing:  
tensor([[3.1482]], requires\_grad=True) Final b:  
Parameter containing:  
tensor([1.9967], requires\_grad=True)

# Linear regression is an old story tough...

... and hence is implemented in *many libraries*. Consider  
Scikit-Learn

-> use the library that they asked.

```
1 from sklearn.linear_model import LinearRegression
2 reg = LinearRegression().fit(X_train.to('cpu').numpy(),
3 y_train.to('cpu').numpy())
4 print("Final w:", reg.coef_, "Final b:\n", reg.intercept_)
```

```
1 Final w: [[3.1484804]] Final b:
2 [1.9966233]
```

## Some concluding remarks & general hints (from my experience)

- If a dedicated software/program does something very specific, it probably is better than a general-purpose program in doing that thing
- **However** from an educational perspective, good to test highly complicated routines on simple examples
- This lecture is fundamental for **Assignment 1**

# Assignment 1 Presentation

....